AI and Games

Todd Davies

July 26, 2016

Overview

The aim of this semester is to understand the following topics.

What are Stackleberg (Leader-Follower) games? What constitutes a solution to a Stackleberg game? How can learning and optimisation be used to learn good solutions? Applications to price setting and marketing will be discussed.

What is reinforcement learning and how is it applied to on-line learning? What are the important mechanisms of on-line learning? How can reinforcement learning be applied to games situations?

Attribution

These notes are based off of both the course notes (¡course notes¿). Thanks to ¡names¿ for such a good course! If you

find any errors, then I'd love to hear about them!

Contribution

Pull requests are very welcome: `https://github.com/Todd-Dav:`
`third-year-notes`

Contents

In Semester One, of this course we covered zero and non-zero sum games (Nash games) and search algorithms to find equlibriums (minimax and alpha beta pruning). In this semester, we are no longer assuming that each player has the same 'role' and that the games are zero-sum. We are letting there be an infinite number of positions for each player to take and not assuming that players have perfect information.

We're going to be looking at Stackelberg games; an example of which could be retail pricing. Different firms might offer different prices on products, they can choose a price from an infinite set (the real numbers) and some firms might have private information that others might not[1].

An industry can be controlled or dominated by a one, two or multiple entities:

Monopoly: When an industry consists of a single firm.

Duopoly: When an industry consists of two firms.

Oligopoly: When an industry consists of a few firms (and when each decision one takes impacts on the other's profits).

In duopoly/oligopoly situations, the roles of the players are different. If one firm chooses to change its prices first, then the others have to decide what to do in light of the change; and may not know the full economic and/or political motivations behind the initial move.

In a two player Stackelberg game, one player selects his strategy first, and then the other responds. The first player

---

[1]Typical examples of unknown information include payoff functions or unit cost.

is called the leader (P_L) and the second is called the follower (P_F). This contrasts with Nash games, where all of the players select their strategies simultaneously.

Both the leader and follower have payoff functions[2], which they want to maximise:

$$J_L(U_L, U_F)$$
$$J_F(U_L, U_F)$$

A strategy space is the set of all possible strategies for a single player. The leader's strategy space is U_L and the follower's is U_F. They can include finite (discrete) or infinite (continuous) strategies. To play the game, the leader must first choose some strategy $u_L \in U_L$ and then the follower must choose some strategy $u_F \in U_F$.

In a Stackelberg game, when the leader has announced u_L, the follower selects a strategy using their reaction function $R(u_L) \in U_F$ such that:

$$J_F(u_L, R(u_L)) = MAX_{u_F \in U_F} J_F(u_L, u_F)$$

This looks complicated, but what it says is that the payoff for the follower is given the parameters u_L and u_F. The follower obviously wants to maximise their payoff, and so tries to find the strategy $u_F \in U_F$ that will maximise their payoff. This isn't always easy, since enumerating every possible

---

[2]A payoff function indicates the benefit for that player given the chosen strategies of all the players.

strategy and knowing what will happen if any given strategy is deployed is sometimes impossible.

In order to solve a Stackelberg game, we need to solve the following problem:

- What is the follower's reaction function $R(u_L)$?

- When we've found that, what leader strategy u_L maximises the leader's payoff function:

$$J_L(u_L^*, R(u_L^*)) = MAX_{u_L \in U_L} J_L(u_L, u_F)$$

If the follower has a reaction function $R(u_L)$ and there exists a leader strategy $u_L^* \in U_L$ and the response strategy is $u_F^* = R(U_L^*) \in U_F$ such that $J_L(u_L^*, u_F^*) = MAX_{u_L \in U_L} J_L(u_L, R(u_L))$, then (u_L^*, u_F^*) is called a **Stackelberg strategy/equilibrium**; i.e. both sides are paying an optimal strategy.

We always assume that the follower is rational and tries to find the best reaction strategy. Sometimes this is untrue, for example if taking a non-optimal strategy results in a rival player having a big loss.

For a player to be a leader, he needs to act first. There is no requirement that they are a leader in an economic or political sense. The only requirement to be a follower is to react to the leader's strategy[3]. As mentioned before, Stackelberg games can be continuous or discrete, depending on if the strategy space is finite or infinite.

Since the only difference between a Nash game and a Stackelberg game is whether moves are played simultaneously or

---

[3]Note that this does not mean the follower is in a weaker position.

in order, should we prefer Nash or Stackelberg games? If the player has the opportunity to be the leader, then Stackelberg games should be preferred, since he will always be better off than in a Nash game in this instance. Here's a mini-proof:

Let u_1 and u_2 be the strategies for player's one and two, and let $J_1(u_1, u_2)$, $J_2(u_1, u_2)$ be their respective payoff functions. If there is a Stackelberg strategy and a Nash strategy, then:

$$J_1(u_1^{Stackelberg}, u_2^{Stackelberg}) \geq J_1(u_1^{Nash}, u_2^{Nash})$$

Note that sometimes the follower can be better off in a Stackelberg game, but the leader would still be better off playing the Stackelberg game than playing a Nash game. Other times, the player could win as the leader or follower, but win better as the follower.

For the follower to play the game, he needs to know nothing about the leader or his strategy space. However, for the leader to play, he needs to know the follower's payoff function and strategy space in order to work out the action that will give the max payoff. If such information is not available, then the leader must guess or learn it (based off previous experience or data).

1 Solving Stackelberg game problems

We can solve with two sequential maximisation problems, each for a single player:

1. For each of the leader's strategies $u_L \in U_L$, solve:

$$max_{u_F \in U_F}(J_F(u_L, u_F)$$

   The solution for this is the follower's reaction function $R(u_L)$.

2. Find the best strategy for the leader by solving:

$$max_{u_L \in U_L}(J_L[u_L, R(u_L)])$$

   If u_L^* is the strategy found in the above maximisation, and $u_L^* = R(u_L^*)$, then (u_L^*, u_F^*) is a Stackelberg strategy.

So, in order to find the Stackelberg strategy, we need to solve two maximisation problems. This is A-level maths, but we can recap it in the next bit.

1.1 How to find the maximum point of a function over a continuous space

We want to find a point $x^* \in X$ that maximises $f(x)$ over X, such that $\forall x \in X, f(x^*) \geq f(x)$. Such a point is called the global maximum point of $f(x)$ on X. A local maximum point y^* is the maximum value of $f(x)$ within the region X, such that $\forall x \in X, f(y^*) \geq f(x)$.

Figure 1: Wikipedia's pictorial explanation of minima and maxima.

The standard method for trying to find a maximum of a function is by using derivatives. A derivative is usually notated by a dash after the function. Finding the derivative of a derivative is called the second order derivative function, and has two dashes.

A first order derivative represents the gradient of the line drawn by the function at a specific point. If $f'(x) > 0$, then the line slopes upwards, and if $f'(x) < 0$ then it's sloping down. Obviously if it is equal to zero, then the line is horizontal.

1.1.1 Calculating a derivative

The most simple rule is:

$$f(x) = x^n \therefore f'(x) = nx^{n-1}$$

This means if the function is constant, it disappears:

$$f(x) = C \therefore f'(x) = 0$$

Here are three simple examples:

$$f(x) = x^2 \therefore f'(x) = 2x$$

$$f(x) = x^3 \therefore f'(x) = 3x^2$$

$$f(x) = 5 \therefore f'(x) = 0$$

If a function is composed of other functions that are added together (of the form $f(x) = f_1(x) + \cdots + f_n(x)$ then:

$$f'(x) = f_1'(x) + \cdots + f_n'(x)$$

Just like:

$$f(x) = -2x^2 + 4x + 5 \therefore f'(x) = -4x + 4$$

Finally, if it's composed of functions that are multiplied together, then $f'(x) = f_1'(x)f_2(x) + f_1(x)f_2'(x)$:

$$f(x) = x^2 \times x \therefore f'(x) = [x^2]' \times x + x^2 \times [x]' = 2x \times x + x^2 \times 1 = 2x^2 +$$

1.1.2 Finding maxima

Once you've found the derivative of your function, and then found the points where the derivatives are 0 (and the gradient of the line is therefore zero), you need to determine whether that point is a minima or maxima. To do this, we differentiate again to get the second order derivative. If the value of $f''(x) \geq 0$ then the gradient is increasing, therefore it's a minima. Because of this, we want points where the second order derivative gives a negative value, indicating that the point is a maxima.

To find the largest maxima, we simply find the one that is largest value of $f(x)$.

2 Solving Stackelberg game problems

We know that there are two steps to solve a Stackleberg game; first of all, you must solve the maximisation problem $\max u_F \in U_F J_F(u_L, u_F)$, giving the reaction function $R(u_L)$. Then you must find the best leader strategy by solving $\max_{u_L \in U_L} J_L(u_L, R(u_L))$ such that (u_L, u_F) is the Stackelberg strategy.

The strategy space (i.e. the different prices that the leader and follower can set) is $U_L = [c_L, +\infty], U_F = [c_F, +\infty]$, where c_L, c_F are the cost of each unit for the leader and follower respectively.

The payoff function (i.e. how much profit the leader and follower make) is defined by:

$$J_L(u_L, u_F) = (u_L - c_L) \times S_L(u_L, u_F)$$

$$J_F(u_L, u_F) = (u_F - c_F) \times S_F(u_L, u_F)$$

This is essentially the profit per unit (sale price minus cost) multiplied by the number of units sold in the sale. The sale function is given to you in exam questions and is a quadratic function over u_L and u_F.

We want to find the optimal values for u_L and u_F. To do this, we first differentiate J_F., then we differentiate J_L and sub in the value for u_F we found in the first stage so we can get answers. Lets do an example; given the following situation,

we want to find the follower's reaction function, and then an optimal values for u_L to give us maximum profit:

$$J_L(u_L, u_F) = (u_L - c_L) \times S_L(u_L, u_F)$$

$$J_F(u_L, u_F) = (u_F - c_F) \times S_F(u_L, u_F)$$

$$c_F = 1 = c_F$$

$$S_L = 5 - 2u_L + u_F$$

$$S_F = 6 + u_L - 2u_F$$

To get the reaction function, we differentiate J_F:

$$
\begin{aligned}
0 &= \frac{d}{du_F} J_F(u_L, u_F) \\
&= \frac{d}{du_F}(u_F - 1)(6 + u_L - 2u_F) \\
&= \frac{d}{du_F} 6u_F + u_F u_L - 2u_F^2 - 6 - u_L + 2u_F \\
&= u_L - 4u_F + 8
\end{aligned}
$$

If we differentiate again, we see that this is a maxima ($-4 > 0$).

We can rearrange it so that: $J'_f(U_L, U_F) = R(u_L) = \frac{u_L}{4} + 2$.

Now we can sub in the value for u_F into J_L:

$$J_L(u_L, R(u_L)) = (u_L - 1)(502u_L + R(u_L))$$
$$= \frac{1}{4}(35u_L - 7u_L^2 - 28)$$

And differentiate:

$$\frac{d}{du_L} J_L(u_L, R(u_L)) = \frac{1}{4}(35 - 14u_L) = 0$$

If we rearrange that, we get $u_L = \frac{35}{14}$, therefore $u_F = 2.625$, which are the optimum values of U_L and U_F in terms of the leader profit.

This means the profits are:

$$J_L(\frac{35}{14}, 2.625) = (\frac{35}{14} - 1)(5 - 2\frac{35}{14} + 2.625) = 3.9375$$

$$J_F(\frac{35}{14}, 2.625) = (2.625 - 1)(6 + \frac{35}{14} - 2(2.625)) = 5.28125$$

3 Stackelberg games with imperfect information

It is very rare that both players in a two person Stackelberg game will have perfect information; in any case, it is in their interest to hide information from one-another. In a game of imperfect information, the follower is not affected; their job is the same (respond to whatever the leader does).

However, this is a huge change for the leader. Without knowing the strategy space for the follower, or the follower's reaction function, they can't work out a good strategy. We need to come up with a technique to let the leader 'guess' a good strategy based of what we *do* know about the game.

There are two obvious solutions, and one machine learning solution to this. We're interested in the ML solution, but the obvious solutions are:

- Choose to be the follower. Then we don't have to decide on the first strategy, and can just choose a reaction. However, this isn't always possible (sometimes you're forced to move first).

- Find the best strategy for the worst scenario; that is to say we should find the best, worst strategy:

$$u_L = argmax_{u_L \in U_L}(min_{u_F \in U_F} J_L(u_L, u_F))$$

Though this gives a lower bound on the payoff, it is also far too conservative. Imagine if you always planned for the worst and never considered (never mind hoped) for the best!

However, the best solution is to attempt to learn what the follower will do based on what has happened in the past. Many two person Stackelberg games are played repeatedly (e.g. setting oil prices every day), and therefore there is a lot of information available to learn from. All we need to learn is the follower's reaction function; i.e. what value they will choose for u_F given any chosen u_L.

The approach we're going to use to do this is called linear regression. This involves finding a linear function that gives similar values to what the follower has done in the past. You could also use a polynomial function, a neural network (neural networks are called the *universal approximation*, since they can mimic any function), or many other different techniques.

So, if $y = R(x)$ be the unknown function, and imagine we have a set of T datapoints where we know the input x and the value y. We want to make a new function $\hat{R}(x)$ that will approximate $R(x)$.

We define $\epsilon(x) = R(x) - \hat{R}(x)$, and try to make ϵ as small as possible for all values of x. Unfortunately, since we don't know R, we can't work out ϵ. However, we can work out the value of ϵ for all the datapoints we do have. Making the following function (called the least square criterion) as small as possible is the aim of the regression:

$$\sum_{t=1}^{T} \epsilon^2[x(t)] = \sum_{t=1}^{T} \left(y(t) - \hat{R}(x(t)) \right)^2$$

Designing a function is problem dependent, but we know our

reaction function will be linear (because we're not considering other functions in the course):

$$R(x) = a + bx$$

To make $\hat{R}(x)$ imitate $R(x)$, we must find values for a and b that minimise the squared criterion we had before:

$$min_{a,b} \sum_{t=1}^{T} (y(t) - (a + bx(t))))^2$$

Instead of solving the minimising problem, we can instead solve the following maximisation problem:

$$max_{a,b} - \sum_{t=1}^{T} (y(t) - (a + bx(t))))^2$$

We can use partial differentiation to do this (just like in `COMP36212`), but it's tricky, so this is what you will eventually get:

$$a = \frac{\sum_{t=1}^{T} x^2(t) \sum_{t=1}^{T} y(t) - \sum_{t=1}^{T} x(t) \sum_{t=1}^{T} x(t)y(t)}{T \sum_{t=1}^{T} x^2(t) - (\sum_{t}^{T} = 1x(t))^2}$$

$$b = \frac{T \sum_{t=1}^{T} x(t)y(t) - \sum_{t=1}^{T} x(t) \sum_{t=1}^{T} y(t)}{T \sum_{t=1}^{T} x^2(t) - (\sum_{t}^{T} = 1x(t))^2}$$

4 Multivariable regression

Now we know how to solve a linear regression problem with one variable, but often, there isn't just one variable. If there are three fuels in a petrol station (diesel, unleaded, super unleaded), then there will be three variables, and we want to learn strategies for all of them.

Here, the leader & follower strategies will be:

$$u_L = (u_1^L, u_2^L, u_3^L)$$

$$u_R = (u_1^R, u_2^R, u_3^R)$$

The follower's reaction function is:

$$\hat{R}(u_L) = \hat{A} + \hat{B}u_L \begin{bmatrix} \hat{a}_1 \\ \hat{a}_2 \\ \dots \\ \hat{a}_f \end{bmatrix} + \begin{bmatrix} \hat{b}_{1,1} & \hat{b}_{1,2} & \dots & \hat{b}_{2,n_L} \\ \hat{b}_{2,2} & \hat{b}_{2,2} & \dots & \hat{b}_{2,n_L} \\ \dots & \dots & \dots & \dots \\ \hat{b}_{n_f,1} & \hat{b}_{n_f,2} & \dots & \hat{b}_{n_f,n_L} \end{bmatrix} \begin{bmatrix} \hat{u}_1^L \\ \hat{u}_2^L \\ \dots \\ \hat{u}_{n_L}^L \end{bmatrix}$$

Given the leader's strategy (u_L), the follower will react using:

$$\hat{R}_i(u_L) = \hat{a}_i + \sum_{j=1}^{n_L} \hat{b}_{i,j} u_j^L$$

In essence, we're trying to find the mathematical relationship between an output variable and several input variables. All

of them take continuous values (if the dependent is discrete, then it's a classification problem). Such problems are widely used in prediction and forecasting.

The idea is that we want to find values for θ that we can sub into $\hat{R}$, such that:

$$\theta^* = arg\ min_\theta \sum_{t=1}^{T} \epsilon^2(t)$$

$$= arg\ min_\theta \sum_{t=1}^{T} (R[X(t)] - \hat{R}[X(t), \theta])^2$$

$$= arg\ min_\theta \sum_{t=1}^{T} (y(t) - \hat{R}[X(t), \theta])^2$$

Since $\hat{R}(X, \theta) = (1, x_1 \ldots, x_n) \begin{pmatrix} a_0 \\ a_1 \\ \vdots \\ a_n \end{pmatrix} = a_0 + a_1 x_1 + \cdots + a_n x_n$, we just need to find a solution to:

$$\theta^* = arg\ min_\theta \sum_{t=1}^{T} (y(t) - [a_0 + a_1 x_1(t) + \cdots + a_n x_n(t)])^2$$

Where $\theta = (a_0, a_1, \ldots, a_n)$.

5   Learning in games

There are two types of Stackelberg strategy generators; online and offline ones. In offline learning, the reaction function

is regarded as a linear regression problem, and is solved by applying the least square method to the historical data. All the learning is done before the playing starts.

In online learning, we aim to carry on learning even while the game is being played. The reaction strategy from the follower is updated as we play the game. We need to work out how to incorporate the follower's moves into our strategy.

The reason why online methods are preferred, is that the environment in which the game is played is always changing; costs may change, share prices may change etc. The follower's reaction function will change over time to reflect this.

There are two commonly used methods:

Moving Window Approach:
  Here, we only consider n historical points, and as new data comes in, we discard old data. To make it better, we can add a parameter λ to make old datapoints mean less than new ones. The impact of the nth most recent datapoint is $n \times \lambda^{n-1}$, and $\lambda \approx 0.95 - 0.99$.

Recursive Least Square Approach:

Social surplus is essentially a measure of if the people who wanted the items most got the items in a bid. For example, if a musician who really needed a car could only bid £500 (though he may value the car at more than this), but a millionaire wanted to buy lots of cars for his children to rally race in and outbid him, despite valuing the car at only £501, then the social surplus would be low.

As we've seen, the design problem for single item auctions is to make a system where one item is sold, n bidders submit a bit for the item (who all want to maximise their utility; value - price paid). We want to solicit the bids and choose the winner, while hopefully maximising the **social surplus**, i.e. the value of the winner.

As we've seen, charging the winner the price of the second highest bid is the Vickrey auction, and is widely used. It's pretty good because it has strong incentives (DSIC), maximises the social surplus and is computationally efficient.

However, it's also fairly constricted in what it can do; we want a more general solution.

Single parameter auction games are where:

- k items for sale and each item consists of several units of a product or service $I_k = \{a_j\}, j = 1, \ldots, k$.

- n bidders, and bidder i has his private value per-unit v_i of the product (called his unit valuation).

- The objective of bidder i is to maximise his quasilinear utility:

$$u_i = v_i \cdot x_i - p_i$$

  - x_i is the number of the product gained if bidder i wins the item, 0 otherwise.

  - p_i is the price paid if bidder i wins, 0 otherwise.

Situation for the seller (mechanism designer) is as follows:

- The outcome space is $x = (x_1, \ldots, x_n)$, where x_i is the number of units (for a given item) being allocated to bidder i.

- The seller wants to maximise social surplus F:

$$F(x_1, \ldots, x_n) = \sum_{i=1}^{n} v_i x_i$$

- In a single parameter environment, there is a constraint meaning that one agent can be allocated at most one item(with x_i units in it).

The problem is that we want to design a *sealed bid auction*, and need to come up with implementations for the following three steps:

- Collect the bids $b = (b_1, \ldots, b_n)$

- Allocate the number of units given to each bidder:

$$x(b) = (x_1(b), \ldots, x_n(b)) \in X \subset R^n$$

- Choose how much each bider must pay:

$$p(b) = (p_1(b), \ldots, p_n(b)) \subset R^n$$

5.0.3 Case study: Sponsored search advertising auctions

In an auction selling search advertising, we want the mechanism to exhibit the following characteristics:

- DSIC; Truthful bidding should always be the dominant strategy and never lead to a negative utility.

- We want to maximise $\sum_{i=1}^{n} v_i x_i$ (social surplus).

- Each bidder gets only one item, and each item can only be assigned to one bidder. This can be defined mathematically:

  There are k items $(k > 0)$, and n bidders. $x = (x_1, \ldots, x_n)$ is the assignment of items to bidders, and x_i is either 0 or 1, representing whether bidder i won an item. The following must hold:

  $$\sum_{i=1}^{n} x_i \leq k$$

- We want the algorithm to run in polynomial time.

Each advertising slot has a Click Through Rate (CTR), which represents how often people who see the ad click on the ad. The higher the CTR for a slot, the more desirable it is. There are:

- k slots for sale, and slot j includes α_j units, which is the CTR for that slot.

- n bidders, where each bidder i has his private value v_i, representing the value he gives to a single click. Thus, his valuation of a slot j is $v_i\alpha_j$.

The bidder's objective is to maximise his own utility of course. Each bidder i will do this by maximising:

$$u_i = v_i \cdot x_i - p_i$$

Essentially the utility is the value he assigns to each click multiplied by the number of clicks he will get, minus the price the bidder will pay if he gets the slot.

Remember, we're trying to find an algorithm similar to the Vickery method, one that is DSIC, maximises social surplus and has a polynomial running time. To design such an algorithm we can work in two steps:

Step 1:
  Assume that the bidders will bid truthfully (assume DSIC). Decide how to assign bidders to items as to maximise social surplus with a poly time algorithm.

Step 2:
  Given our answer for step 1, decide how to set our selling prices so that DSIC does indeed hold.

For step one, if we assume that we've got more bidders than items $(n \geq k)$, and we know the value each bidder assigns to one click (v_i), then we can maximise social surplus by

allocating the ith largest item to the bidder with the ith highest valuation.

To do this, we can sort the valuations $v_1 \geq \cdots \geq v_x$ in polynomial time, and then sort the items in order of CTR (again, polynomial). Then we can easily get an assignment x^* since we simply match the bidders' valuations against the items; the bidder at valuation index i gets the item at index i. If there are more bidders than items, then bidders at and index greater than the number of items get 0.

As the above allocation rule is a function of the bid (or the valuation of the CTR), then it is a **Monotone Allocation Rule**, meaning that bidding higher can only get you more stuff.

Answering the step two is now possible since we have an allocation rule. We now need to determine how to set prices such that DSIC will hold. First we need to define an **Implementable Allocation Rule**:

> An allocation rule $x(b)$ for a single parameter environment is implementable if there is a payment rule $p(b)$ that makes the sealed-bid auction $[x(b), p(b)]$ DSIC.

If an allocation rule is implementable, this means that you can *always* define a payment rule that makes truthful bidding the dominant strategy for each bidder. This is **Myerson's Lemma**:

- An allocation rule is implementable iff it is monotone.

- If the allocation rule is monotone, then there is a payment rule that makes the sealed bid mechanism DSIC.

- This payment rule can be given by an explicit formula:
  - Assume that the items for sale are sorted by the number of units they represent.
  - Also assume that the bids are sorted.
  - The payment rule is therefore:

$$p_i(b) = p_i(b_1, \ldots, b_i, \ldots, b_k, \ldots, b_n)$$

$$= \sum_{j=1}^{k} b_{j+1}(\alpha_j - \alpha_j + 1)$$

This pricing rule is DSIC because if the bidder sets $b_i = v_i$:

$$u_i(v_i, b_i) = v_i \times \alpha_i - p_i(b)$$

This is basically that the utility of the bidder is the value they assign to on click multiplied by the CTR they won, minus the price for that bid. It can be calculated that bidding lower or higher than v_i gives a lower value for the utility function, since if $b_i < v_i$ then the bidder may get a smaller item and lose profitable units, and if $b_i > v_i$ then the bidder may get a bigger item, but these additional units are not profitable. =======

6   Mechanism Design

Mechanism design involves taking an engineering approach to designing economic mechanisms or incentives in order to

steer the outcome of events/transactions/games etc towards a desired objective. As with the Nash games in semester one, and the Stackleberg games in the first part of semester two, it is always assumed that the players will act rationally. Since we start from the desired outcomes, and design a game to produce these outcomes, mechanism design is also known as *reverse game theory*.

6.1   A toy example

A simple example is that of a mother dividing a cake for her two children, with a view to them both getting exactly half since otherwise they will complain. If the mother cuts the cake as equally as she can, then there is a risk that one or both of the children will have a different view on the size of the pieces to her, and will want a specific piece.

> **Note:**
>
> Think about how this would work for n children.

The solution is to have one child cut the cake into two pieces, but then let the other child have first pick of the pieces. Then, no matter if the first child cuts the cake evenly or not, the second child will take whichever one that they think is biggest, and since the first child cut the cake in the first place, they shouldn't (in theory) contest the choice.

6.2   A real example

There are lots of more interesting examples of mechanism design, such as the government assigning licences to mobile operators. If we view the government as one player, and the different mobile operators as n different players, then it is easy to see their differing aims. The government wants to allocate a licence to whoever values it the most highly since an operator who highly values the licence is likely to use it best for their customers and maximise *social surplus*. The goal of mobile operators is to maximise their own profit (i.e. to minimise the price they actually pay for the licence).

If the government could simply ask for an honest recommendation of each mobile operator's valuation of the licence, then they could simply give the licence to the highest valuing operator. However, since the government does not know the *true* valuation (as opposed to whatever the operator may tell the government) of the licence by each operator, they cannot really know who to give it to.

Other than asking for the valuation, the government could also require the operators to pay their valuation when they submit it (so that they don't overstate their valuation), but then they would likely understate their valuation to save money. The solution is to allocate the licence to the operator with the highest valuation, but have them pay the *second highest* valuation.

If this is done, then the operator would not report a higher valuation than their true valuation as it would not be profitable (since they would only pay the second highest), if they bid lower than their true valuation they might lose the

chance of winning, and the operators would be willing to participate in the first place (and bid their true valuation) since it is profitable if the win, and they don't lose if they don't win.

6.3 Defining 'Mechanism Design'

From the lecture notes:

> Mechanism design is the sub-field of microeconomics and game theory that considers how to implement good system-wide solutions to problems that involve multiple self-interested agents, each with private information about their preferences.

Overall, there are n players (aka agents), and each player has some private information and will try to maximise their own payoff (aka utility) function[4].

There is a mechanism designer such as the government in the mobile operator's spectrum example above. The designer does not know the private information of each agent, but wants to achieve some desired outcome such as maximising the social welfare (e.g. a government) or their own profit (e.g. a private seller).

So, we need to design the rules of the game in order to coerce a desirable outcome based on agent strategies.

---

[4]Note that the strategy space over which the payoff function is defined is not itself defined until the mechanism has been designed.

6.3.1 A Mathematical Definition

The outcome is the set Ω, and the players are $i \in I$ where I is the set of players and $|I| = N$. The private information of each player is $\theta_i \in Theta_i$.

The utility function for each player is $u(o, \theta_i)$ for the player's information set θ_i and the outcome $o \in \Omega$.

A mechanism is a pair $M = (S, g)$ where S is the strategy space $S^N = S_1 \times \cdots \times S_N$ where player I chooses a strategy $s_i(\theta_i) \in S_i$ and s_i maps the set of information sets onto the set of strategies available to the player $s_i : \Theta_i \to S_i$. The outcome function is $g : S^N \to \Omega$.

Finally, a game is the utility to player i from their strategy profile; $u_i(g(s(\theta)), \theta_i)$.

We want to make a mechanism $M = (S, g)$ where $g(s_1(\theta_1)), \ldots, s_N$ $f(\theta), \forall \theta \in \Theta^N$, such that $(s_1, \ldots, s_N)$ is an equilibria.

6.4 Implementing Mechanisms

There is always a danger that if there are multiple equilibria in a mechanism, then the participants of the game may end up stuck in some sub-optimal equilibria. We want to create a mechanism where all of its equilibria are optimal to avoid this conundrum, which can only be done if the mechanism is a maskin monotonicity.

6.5 Why Do We Care About Mechanism Design?

Mechanism design has wide applications in applied CS, such as in internet routing, job scheduling, peer to peer systems, cryptography etc. Mechanism design has some complicated computational problems, and often requires effective algorithms to solve it.

6.6 Single Item Auctions

If we want to sell one item to n bidders who all want to maximise their utility (the value they give the item minus the price they pay for it), then we need to create an auction to solicit the bids and choose the winner (and how much the winner pays). At first, this seems like a trivial problem, but in fact it has many intricacies.