

AI and Games

Todd Davies

February 4, 2016

Overview

The aim of this semester is to understand the following topics.

What are Stackleberg (Leader-Follower) games? What constitutes a solution to a Stackleberg game? How can learning and optimisation be used to learn good solutions? Applications to price setting and marketing will be discussed.

What is reinforcement learning and how is it applied to on-line learning? What are the important mechanisms of on-line learning? How can reinforcement learning be applied to games situations?

Attribution

These notes are based off of both the course notes (course notes). Thanks to names for such a good course! If you

find any errors, then I'd love to hear about them!

Contribution

Pull requests are very welcome: <https://github.com/Todd-Davis/third-year-notes>

Contents

- 1 Solving Stackleberg game problems
 - 1.1 How to find the maximum point of a function over a continuous space
 - 1.1.1 Calculating a derivative

In Semester One, of this course we covered zero and non-zero sum games (Nash games) and search algorithms to find equilibriums (minimax and alpha beta pruning). In this semester, we are no longer assuming that each player has the same ‘role’ and the games are zero-sum. We are letting there be an infinite number of positions for each player and not assuming perfect information.

We’re going to be looking at Stackleberg games; an example of which could be retail pricing. Different firms might offer different prices on products, they can choose infinite price points (or close enough) and some firms might have private information that others might not.

An industry can be controlled or dominated by a one, more or entity:

Monopoly: When an industry consists of a single firm.

Duopoly: When an industry consists of two firms.

Oligopoly: When an industry consists of a few firms (and when each decision one takes impacts on the other’s profits).

In duopoly/oligopoly situations, the roles of the players are different. If one firm chooses to change its prices first, then the others have to decide what to do in light of the change; and may not know the full economic and/or political motivations behind the initial move.

In a two player Stackleberg game, one player selects his strategy first, and then the other responds. The first player is called the leader (P_L) and the second is called the follower (P_F). In a Nash game, the players select their strategy

simultaneously.

There are payoff functions for the leader and follower, where each player wants to maximise their own:

$$J_L(U_L, U_F)$$

$$J_F(U_L, U_F)$$

A strategy space is the set of all possible strategies for a single player. The leader's strategy space is U_L and the follower's is U_F . They can include finite (discrete) or infinite (continuous) strategies. The leader must choose $u_L \in U_L$ and the follower must choose $u_F \in U_F$.

In a Stackleberg game, the leader first announces his strategy u_L , and then the follower selects his best response $R(u_L) \in U_F$ which maximises:

$$J_F(u_L, R(u_L)) = \text{MAX}_{u_F \in U_F} J_F(u_L, u_F)$$

In order to solve a Stackleberg game, we need to solve the following problem:

- What is the follower's reaction function $R(u_L)$?
- When we've found that, what leader strategy u_L maximises the leader's payoff function:

$$J_L(u_L^*, R(u_L^*)) = \text{MAX}_{u_L \in U_L} J_L(u_L, u_L)$$

If the follower has a reaction function $R(U_L)$ and there exists a leader strategy $U_L^* \in U_L$ and the response strategy is $u_F^* = R(U_L^*) \in U_F$ such that $J_L(u_L^*, u_F^*) = \text{MAX}_{u_L \in U_L} J_L(u_L, R(u_L))$, then (u_L^*, u_F^*) is called a **Stackleberg strategy/equilibrium**.

We always assume that the follower is rational and tries to find the best reaction strategy. Sometimes this is untrue, for example if taking a non-optimal strategy results in a rival player having a big loss.

For a player to be a leader, he needs to act first. There is no requirement that they are a leader in an economic or political sense. The only requirement to be a follower is to act second; the follower is not necessarily in a weaker position. As mentioned before, Stackleberg games can be continuous or discrete, depending on if the strategy space is finite or infinite.

Since the only difference between a Nash game and a Stackleberg game is whether moves are played simultaneously or in order, should we prefer Nash or Stackleberg games? If the player has the opportunity to be the leader, then Stackleberg games should be preferred, since he will always be better off than in a Nash game in this instance. Here's a mini-proof:

Let u_1 and u_2 be the strategies for player's one and two, and let $J_1(u_1, u_2)$, $J_2(u_1, u_2)$ be their respective payoff functions. If there is a Stackleberg strategy and a Nash strategy, then:

$$J_1(u_1^{Stackleberg}, u_2^{Stackleberg}) \geq J_1(u_1^{Nash}, u_2^{Nash})$$

I don't understand the next bit on page 29...

Note that sometimes the follower can be better off in a Stackleberg game, but the leader would still be better off playing the Stackleberg game than playing a Nash game. Other times, the player could win as the leader or follower, but win better as the follower.

For the follower to play the game, he needs to know nothing about the leader or his strategy space. However, for the leader to play, he needs to know the follower's payoff function and strategy space in order to work out the action that will give the max payoff. If such information is not available, then the leader must guess or learn it (based off previous experience or data).

1 Solving Stackleberg game problems

We can solve with two sequential maximisation problems, each for a single player:

1. ...
2. ...

1.1 How to find the maximum point of a function over a continuous space

We want to find a point $x^* \in X$ that maximises $f(x)$ over X . In other words $\forall x \in X, f(x^*) \geq f(x)$. Such a point is a global maximum point of $f(x)$ on X . A local maximum point y^* is the maximum in the neighbourhood U of $y^* \in X$.

1.1.1 Calculating a derivative

Since the derivative is the gradient of a function at a point:

Note:

See my A-level math notes maybe!

$$f(x) = C \therefore f'(x) = 0$$

$$f(x) = x \therefore f'(x) = 1$$

$$f(x) = x^2 \therefore f'(x) = 2x$$

$$f(x) = x^3 \therefore f'(x) = 3x^2$$

$$f(x) = -2x^2 + 4x + 5 \therefore f'(x) = -4x + 4$$

$$f(x) = x^2 \times x \therefore f'(x) = [x^2]' \times x + x^2 \times [x]' = 2x \times x + x^2 \times 1 = 2x^2 + x$$

Once you've found the points where the gradients are 0 using the first order derivatives, then you need to differentiate again and find if it's a minimum or maximum.