

AI and Games

Todd Davies

February 4, 2016

Overview

The aim of this semester is to understand the following topics.

What are Stackleberg (Leader-Follower) games? What constitutes a solution to a Stackleberg game? How can learning and optimisation be used to learn good solutions? Applications to price setting and marketing will be discussed.

What is reinforcement learning and how is it applied to on-line learning? What are the important mechanisms of on-line learning? How can reinforcement learning be applied to games situations?

Attribution

These notes are based off of both the course notes (jcourse notes_i). Thanks to jnames_i for such a good course! If you find any errors, then I'd love to hear about them!

Contribution

Pull requests are very welcome: <https://github.com/Todd-Davies/third-year-notes>

Contents

1	Solving Stackelberg game problems	4	1.1	How to find the maximum point of a function over a continuous space	4
			1.1.1	Calculating a derivative	5
			1.1.2	Finding maxima	6

In Semester One, of this course we covered zero and non-zero sum games (Nash games) and search algorithms to find equilibriums (minimax and alpha beta pruning). In this semester, we are no longer assuming that each player has the same ‘role’ and the games are zero-sum. We are letting there be an infinite number of positions for each player and not assuming perfect information.

We’re going to be looking at Stackelberg games; an example of which could be retail pricing. Different firms might offer different prices on products, they can choose infinite price points (or close enough) and some firms might have private information that others might not.

An industry can be controlled or dominated by a one, more or entity:

Monopoly: When an industry consists of a single firm.

Duopoly: When an industry consists of two firms.

Oligopoly: When an industry consists of a few firms (and when each decision one takes impacts on the other’s profits).

In duopoly/oligopoly situations, the roles of the players are different. If one firm chooses to change its prices first, then the others have to decide what to do in light of the change; and may not know the full economic and/or political motivations behind the initial move.

In a two player Stackelberg game, one player selects his strategy first, and then the other responds. The first player is called the leader (P_L) and the second is called the follower (P_F). In a Nash game, the players select their strategy simultaneously.

There are payoff functions for the leader and follower, where each player wants to maximise their own:

$$J_L(U_L, U_F)$$

$$J_F(U_L, U_F)$$

A strategy space is the set of all possible strategies for a single player. The leader’s strategy space is U_L and the follower’s is U_F . They can include finite (discrete) or infinite (continuous) strategies. The leader must choose $u_L \in U_L$ and the follower must choose $u_F \in U_F$.

In a Stackelberg game, the leader first announces his strategy u_L , and then the follower selects his best response $R(u_L) \in U_F$ which maximises:

$$J_F(u_L, R(u_L)) = \text{MAX}_{u_F \in U_F} J_F(u_L, u_F)$$

In order to solve a Stackelberg game, we need to solve the following problem:

- What is the follower’s reaction function $R(u_L)$?
- When we’ve found that, what leader strategy u_L maximises the leader’s payoff function:

$$J_L(u_L^*, R(u_L^*)) = \text{MAX}_{u_L \in U_L} J_L(u_L, u_L)$$

If the follower has a reaction function $R(U_L)$ and there exists a leader strategy $U_L^* \in U_L$ and the response strategy is $u_F^* = R(U_L^*) \in U_F$ such that $J_L(u_L^*, u_F^*) = \text{MAX}_{u_L \in U_L} J_L(u_L, R(u_L))$, then (u_L^*, u_F^*) is called a **Stackelberg strategy/equilibrium**.

We always assume that the follower is rational and tries to find the best reaction strategy. Sometimes this is untrue, for example if taking a non-optimal strategy results in a rival player having a big loss.

For a player to be a leader, he needs to act first. There is no requirement that they are a leader in an economic or political sense. The only requirement to be a follower is to act second; the

follower is not necessarily in a weaker position. As mentioned before, Stackelberg games can be continuous or discrete, depending on if the strategy space is finite or infinite.

Since the only difference between a Nash game and a Stackelberg game is whether moves are played simultaneously or in order, should we prefer Nash or Stackelberg games? If the player has the opportunity to be the leader, then Stackelberg games should be preferred, since he will always be better off than in a Nash game in this instance. Here's a mini-proof:

Let u_1 and u_2 be the strategies for player's one and two, and let $J_1(u_1, u_2)$, $J_2(u_1, u_2)$ be their respective payoff functions. If there is a Stackelberg strategy and a Nash strategy, then:

$$J_1(u_1^{Stackelberg}, u_2^{Stackelberg}) \geq J_1(u_1^{Nash}, u_2^{Nash})$$

I don't understand the next bit on page 29...

Note that sometimes the follower can be better off in a Stackelberg game, but the leader would still be better off playing the Stackelberg game than playing a Nash game. Other times, the player could win as the leader or follower, but win better as the follower.

For the follower to play the game, he needs to know nothing about the leader or his strategy space. However, for the leader to play, he needs to know the follower's payoff function and strategy space in order to work out the action that will give the max payoff. If such information is not available, then the leader must guess or learn it (based off previous experience or data).

1 Solving Stackelberg game problems

We can solve with two sequential maximisation problems, each for a single player:

1. For each of the leader's strategies $u_L \in U_L$, solve:

$$\max_{u_F \in U_F} (J_F(u_L, u_F))$$

The solution for this is the follower's reaction function $R(u_L)$.

2. Find the best strategy for the leader by solving:

$$\max_{u_L \in U_L} (J_L[u_L, R(u_L)])$$

If u_L^* is the strategy found in the above maximisation, and $u_L^* = R(u_L^*)$, then (u_L^*, u_F^*) is a Stackelberg strategy.

So, in order to find the Stackelberg strategy, we need to solve two maximisation problems. This is A-level maths, but we can recap it in the next bit.

1.1 How to find the maximum point of a function over a continuous space

We want to find a point $x^* \in X$ that maximises $f(x)$ over X , such that $\forall x \in X, f(x^*) \geq f(x)$. Such a point is called the global maximum point of $f(x)$ on X . A local maximum point y^* is the maximum value of $f(x)$ within the region U , such that $\forall x \in U, f(y^*) \geq f(x)$.

The standard method for trying to find a maximum of a function is by using derivatives. A derivative is usually notated by a dash after the function. Finding the derivative of a derivative is called the second order derivative function, and has two dashes.

A first order derivative represents the gradient of the line drawn by the function at a specific point. If $f'(x) > 0$, then the line slopes upwards, and if $f'(x) < 0$ then it's sloping down. Obviously if it is equal to zero, then the line is horizontal.




Figure 1: Wikipedia's pictorial explanation of minima and maxima.

1.1.1 Calculating a derivative

The most simple rule is:

$$f(x) = x^n \therefore f'(x) = nx^{n-1}$$

This means if the function is constant, it disappears:

$$f(x) = C \therefore f'(x) = 0$$

Here are three simple examples:

$$f(x) = x^2 \therefore f'(x) = 2x$$

$$f(x) = x^3 \therefore f'(x) = 3x^2$$

If a function is composed of other functions that are added together (of the form $f(x) = f_1(x) + \dots + f_n(x)$) then:

$$f'(x) = f'_1(x) + \dots + f'_n(x)$$

Just like:

$$f(x) = -2x^2 + 4x + 5 \therefore f'(x) = -4x + 4$$

Finally, if it's composed of functions that are multiplied together, then $f'(x) = f'_1(x)f_2(x) + f_1(x)f'_2(x)$:

$$f(x) = x^2 \times x \therefore f'(x) = [x^2]' \times x + x^2 \times [x]' = 2x \times x + x^2 \times 1 = 2x^2 + x^2 = 3x^2$$

1.1.2 Finding maxima

Once you've found the derivative of your function, and then found the points where the derivatives are 0 (and the gradient of the line is therefore zero), you need to determine whether that point is a minima or maxima. To do this, we differentiate again to get the second order derivative. If the value of $f''(x) \geq 0$ then the gradient is increasing, therefore it's a minima. Because of this, we want points where the second order derivative gives a negative value, indicating that the point is a maxima.

To find the largest maxima, we simply find the one that is largest value of $f(x)$.